CF Bulk Upload – Demo Brief

The SCB CF bulk upload creates and updates content fragments for six models from one Excel workbook, with a dry run, a results table and a versioned report saved to DAM. Everything below is built, deployed on local AEM (port 4507) and committed; delete, publishing and Phase 2 validation are the open topics.

> This brief describes the tool as demoed on 15 Oct 2026. The feedback from the demo was built on > 20 Oct; see Demo feedback and changes below. The other > documents describe the current behaviour.

Demo feedback and changes (20 Oct)

Feedback	Built
Dry Run button; Final Run enabled only after a dry run	**Dry run** and **Final run** buttons replace the checkbox and Run. Final run is enabled only after a dry run of the same workbook and folder with **no Fail rows**; the server refuses it when the workbook changed after the dry run
Each run saves a report with date and time in its name, in a dedicated folder	Each final run saves `cf-upload-run-<yyyy-MM-dd-HH-mm-ss>.xlsx` (results) in `/content/dam/aemcloudproject/reports`. **Save report to DAM** stays, now saving `offer-listing-report-<yyyy-MM-dd-HH-mm-ss>.xlsx` as a new file each time instead of versions (this also removes the Assets "Replace" risk that wiped the version history)
Users download a sample sheet and author on top of it	**Download sample sheet**: `_Guide` plus every sheet with all column names, no rows
Slug used only for the path; no slug field in the models; existing slug fields stay as they are	A fixed `slug` column sets the path and is not stored. The slug fields were removed from all six models in the repo (decided after the feedback); `nameField` is gone from the profiles. Because the `/conf` filter is `mode="merge"`, instances that already have the models keep the old slug fields until they are removed by hand
Delete option; referenced fragments give a warning and must be fixed first	`DELETE` action. Referenced (other fragments under the picked folder, pages, experience fragments) or published fragments get **Warn** and are not deleted; deletes run last, offers first
Offer details sorted by offer end date, descending	Saved reports sort `offer-detail` by `offerEndDate`, latest first (profile setting `sortField`)
Action dropdown with Create / Update / Delete	Excel dropdown on the action column of the sample and reports
Pathfield for the DAM parent path	**Fragment folder** picker, pre-filled from OSGi (`/content/dam/aemcloudproject/cfs/offer-listing`); profile folders are relative to it (`offers`, `merchants`, …)

Side effect: with an empty sample sheet, new IDs now continue from the highest ID in the sheet **or in DAM**, so they no longer restart at 100001 / 1001.

Demo script

About 15 minutes: template → dry run with errors → real run → update and clear → config.

Prepare before the demo

- ☐ Log in to AEM on 4507 (the local session expires; log in again just before).
- ☐ Click **Save report to DAM** once and download it: with no fragments it is an empty template with `_Guide` plus six header-only sheets.

□ From it, build a demo workbook: one venue, one merchant (pointing to the venue), one category, one card, one CTA and one offer (pointing to them), plus two deliberate mistakes: a slug with a space (`Bad Slug`) and an offer without `offerTitle`.

□ Upload it to `/content/dam/aemcloudproject/imports`.

Ready (15 Oct): two workbooks are already in `/content/dam/aemcloudproject/imports`. `demo-offers.xlsx` gives 6 Pass and 2 Fail (`Festive Bundle` slug, `lunch-special` without `offerTitle`); `demo-offers-fixed.xlsx` is the corrected copy, 8 Pass, as a fallback if you'd rather not edit live. Both only dry-run tested, nothing written.

Show, in this order

1. **Entry point:** Tools → **SCB Tools** → **Content Fragment Bulk Upload**.
2. **Save report to DAM:** open the file; show the `_Guide` sheet (rules, limits, generated table), the red required headers, and the asset's **Timeline** with a version per save.
3. **Dry run** of the demo workbook: green Pass rows with generated IDs (`merchantID 100001`, `offerID 100001`), red Fail rows with clear messages (`offerSlug: "Bad Slug" must be…`, `offerTitle: required`), grey Skip for untouched sheets, **Download CSV**. Point out that nothing was written.
4. Fix the two rows, upload again, **untick Dry run, Run**: fragments appear in Assets, in their folders.
5. **Update and clear:** Save report again (now with rows), set `UPDATE` on the CTA, change its label and empty `deeplink` → dry run shows `(cleared: deeplink)` → real run.
6. **Configuration:** `/system/console/configMgr` → **SCB CF Model Profile** (one per model) and **SCB CF Bulk Upload** (paths, limits, report folder and name). A new model needs one config file, no code.
7. **Authoring Guide** document for authors.

Messages to land: checks before writing; dry run by default; per-row results; config-driven models; no delete or publish by design (yet).

Test scenarios

Every scenario below was run on 15 Oct with the result shown; 4507 was cleaned afterwards. Files are in `/content/dam/aemcloudproject/imports`. **Before:** works on an empty instance. **After:** needs the real run of `demo-offers-fixed.xlsx` first.

#	Scenario	File	When	Expected result
1	Empty template	Save report to DAM	Before	`_Guide` + 6 header-only sheets, red required headers; version in Timeline
2	Create a full set, with mistakes	demo-offers.xlsx (dry run)	Before	6 Pass with IDs (merchantID 100001, ID 1001, offerID 100001); 2 Fail: slug `Festive Bundle`, `offerTitle: required`
3	Create a full set	demo-offers-fixed.xlsx (dry run, then real run)	Before	8 Pass; offers 100001–100003; fragments in their folders
4	Whole workbook rejected	demo-workbook-errors.xlsx	Any	Red alert listing all 3: header with a space, action `UPDTE`, sheet without `action`; nothing checked or written
5	Pick a folder or a non-Excel file	Path field: choose the imports folder itself; for a non-Excel file, upload any image or PDF to /imports first	Any	"… is not a DAM asset" / "… is not an .xlsx file"

#	Scenario	File	When	Expected result
6	Run the same CREATE twice	demo-offers-fixed.xlsx again	After	Every row Fail: "Fragment: already exists at …"
7	Update, clear, partial columns, `_Notes`	demo-updates.xlsx	After	Venue Pass `(cleared: ophr)`; CTA Pass `(cleared: deeplink)`; offer Pass; category Skip; lowercase `update` accepted; `_Notes` ignored
8	Row errors, the rest still runs	demo-row-errors.xlsx	After	1 Pass + 10 Fail: already exists, not found, `offerTitle: required` (emptied on UPDATE), `10001a` not a whole number, `maybe` not yes/no, `01/11/2026` not a date, two CTAs, `#CLEAR` no longer used, `promoCode` not a field, sheet `offers` not a known model
9	Report with data	Save report to DAM	After	Rows per sheet, sorted by ID; new version in Timeline
10	Re-upload the report unchanged	The saved report	After	`_Guide` ignored, every sheet Skip, status ok

Tip for the demo: scenarios 2 → 3 → 7 → 8 tell the story in about 10 minutes. Remember that after scenario 3 the fragments exist, so re-running 2 or 3 gives "already exists".

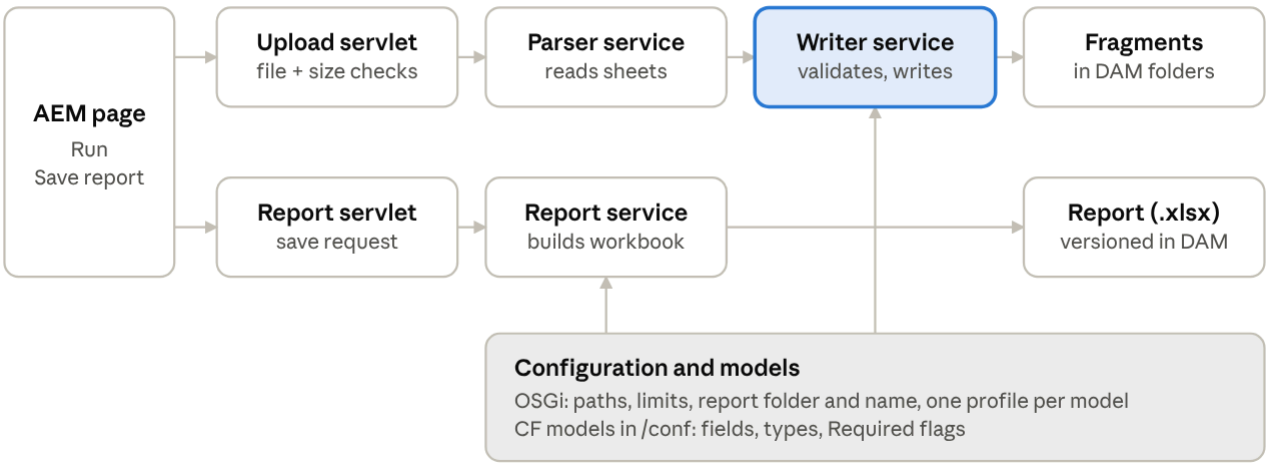
What's built

All of this works end to end on 4507 and is committed (latest commit `3dfd6ef`).

Area	Feature
Workbook	`.xlsx` from DAM; one sheet per model; row 1 = field names; `action` column with CREATE / UPDATE (any case); blank-action rows skipped; sheets starting with `_` ignored
Values	Text, rich text (HTML), whole and decimal numbers, yes/no, dates (with or without time zone), single and multiple fragment references (one path per line), image and link references
Clearing	On UPDATE an empty cell clears the field (except the ID); columns left out are untouched; results list `(cleared: …)`; `#CLEAR` retired and rejected
IDs	Automatic on CREATE: highest ID in the sheet + 1, per model; start values 100001 (offers, merchants) and 1001 (categories, cards); works for number and text IDs; failed rows don't use an ID
Checks before writing	Slug rule; columns must be model fields; required fields (from the model's Required flag); CREATE needs a free path, UPDATE an existing fragment of the same model; values must convert
Safety	Dry run on by default; each row saved on its own (one bad row never undoes others); limits 10 MB and 5,000 rows with an action; workbook-level errors stop everything before any write
Results	Table with sheet, row, slug, action, Pass / Fail / Skip and details; failures first; counts; CSV download
Report	**Save report to DAM**: one sheet per model (header-only when empty, so it's a template), `_Guide` first sheet, required headers in red, rows sorted by ID; saved to a fixed file in a configured folder; the previous file kept as a version (Timeline)
Configuration	OSGi: fragment root, models path, report path, report folder and file name, limits; one **model profile** per model (folder, slug, title, ID field, start ID, sheet order)
Models	Six models aligned with the SCB reference (offer-detail 29 fields, offer-card `ID`, text `merchantID`, single `offerCta`, …)
Entry point	Tools → SCB Tools → Content Fragment Bulk Upload
Docs	Spec in the repo; Authoring Guide for authors; `_Guide` sheet in every report

How it fits together

Upload writes fragments; Save report writes a versioned workbook



Both flows read the same configuration and models, so a new model or a newly required field changes the upload checks and the report together.

Part	Class
Upload servlet	`ScbCfBulkUploadServlet` (POST `upload.json`)
Parser service	`ScbWorkbookParserService` / `ScbWorkbookParserServiceImpl`
Writer service	`ScbContentFragmentWriterService` / `ScbContentFragmentWriterServiceImpl`
Report servlet	`ScbCfReportServlet` (POST `report.json`)
Report service	`ScbContentFragmentReportService` / `ScbContentFragmentReportServiceImpl`
Configuration	`ScbCfBulkUploadConfiguration`; `ScbModelProfileConfiguration` + `ScbModelProfile` / `ScbModelProfileImpl` (one per model)
Shared helpers	`ScbCfBulkUploadUtil` (Required flags, folder creation); `ScbSheetData` (parsed sheet)
Page	Granite page `/apps/aemcloudproject/cfbulkupload/content/cfbulkupload.html` + clientlib `cf-bulk-upload.js` / `.css`

Key decisions taken

Most decisions trade flexibility for safety and simplicity for authors.

Decision	Why
Built from scratch instead of the ACS Commons CF importer	ACS gave poor error reporting, weak field-type handling, timeouts and setup overhead
Excel `.xlsx` picked from DAM, one sheet per model	Familiar to authors; the sheet name is the model name, so no mapping file
Only CREATE and UPDATE; delete stays manual	Delete is irreversible; on hold until the lead decides (see Open decisions)

Decision	Why
Dry run ticked by default, no confirmation popup	Every run is a safe check unless the author deliberately unticks it
All checks before writing; each row saved on its own	A bad row never half-writes a fragment or undoes good rows
Empty action = row skipped; no actions in a sheet = Skip line	A full report can be re-uploaded with only a few rows changed
Empty cell on UPDATE clears the field; `#CLEAR` retired	Lead's proposal: the report holds current values, so deleting a value means "remove it"
IDs: highest in the sheet + 1, start values per model	No AEM lookup needed; works when the sheet contains the existing rows
Required fields come only from the CF model	One source of truth; authoring UI and upload enforce the same rules
Paths, limits and model profiles in OSGi config	A new model or environment needs config, not code
Report saved to DAM (fixed name, previous file versioned) instead of downloaded	Shared location, history in the Timeline; uses `AssetManager` for now
Report is also the template (header-only sheets, `_Guide` first)	Authors always start from a correct workbook; sheets starting with `_` are ignored
No Sling Job (synchronous upload)	Not needed at current volumes; limits guard against timeouts
Models aligned to the SCB reference; existing local fragments deleted	The photos of the SCB models are the source of truth
Concise Javadoc only; `scb` only in class, method, JS and CSS names	Team conventions
Report's slug column = the fragment's node name, not the stored field	The slug is how UPDATE finds a fragment. Legacy fragments with an empty slug field, and fields edited by hand, still round-trip correctly; an UPDATE back-fills or resets the field (tested 15 Oct)

Risks and limitations flagged

The top three for a real rollout: model changes don't deploy to existing environments, there is no undo, and references aren't checked.

#	Risk or limitation	Impact	Mitigation today / proposed
1	`/conf` filter is `mode="merge"`	Model changes in `ui.content` never update models that already exist on stage or prod	Decide the deployment approach for models before going beyond local
2	No undo after a real run	A valid but wrong workbook overwrites values	Dry run, `(cleared: …)` notes, report versions as "before" snapshot; proposed: versions before UPDATE, backup workbook per run
3	References not checked	A path to a missing fragment or image is saved as is; broken links on the frontend	Phase 2 reference check
4	Duplicate IDs possible	IDs come from the sheet only; a workbook without existing rows, or a typed ID, can repeat an ID	Always start from a fresh report; Phase 2 duplicate check (against AEM too)

#	Risk or limitation	Impact	Mitigation today / proposed
5	Duplicate slugs in one workbook	Two CREATE rows both pass in a dry run; two UPDATE rows: the last one wins	Phase 2 duplicate check
6	Header problems surface per row	A wrong column fails every row; a misspelt slug header shows a misleading slug message; a blank header's column is silently ignored	Phase 2: one header check per sheet, clearer slug messages
7	Synchronous upload	Very large uploads can hit the Cloud request time limit (roughly a minute) and lose the results	Limits (10 MB, 5,000 rows); measure ~500 rows and lower the limit if needed; Sling Job if bulk loads become common
8	`AssetManager` is deprecated on Cloud	Saving the report could be restricted later	Works today; direct binary upload is the documented replacement
9	Excel date cells	A real Excel date turns into Excel's display format and fails	Guide tells authors to type dates as text
10	Nothing is published	New and changed fragments stay on author until someone publishes	Manage Publication with "include references"; proposed: publish step in the tool
11	Servlets also deploy to publish	Low risk: the dispatcher blocks `/apps`	Restrict them to author
12	Fragment metadata, variations, language copies	Only master content and the title are set	Out of scope unless needed
13	Multi-value fields that aren't paths	Several values in a text field can't be uploaded	Not needed for the current models
14	Permissions	Authors need write rights on the fragment, imports and report folders	Check groups before rollout

Open decisions

These are waiting on you or the lead; each has a suggestion ready.

Topic	Question	Suggestion
Delete	Decided and built 20 Oct: real delete, referenced or published fragments left in place with Warn, dry-run gate as the safeguard	Open: may a deleted fragment's ID be reused?
Publishing	How to publish exactly what a run wrote	Now: Manage Publication with "include references", checked against the run's CSV. Later: a "Publish N written fragments" step with confirmation
Rollback	Versions before every UPDATE? Backup workbook per real run? Restrict real runs to a group? Show before → after values in the dry run?	Versions before UPDATE + before/after in the dry run first
Duplicates	Fail only the rows involved? Include rows without an action? Check IDs against AEM? (IDs are now generated from sheet + AEM)	Yes to all three
Headers vs model	Bad headers reject the workbook or the sheet? "Did you mean" hint? Check skipped sheets?	Reject the workbook, add the hint, ignore skipped sheets

Topic	Question	Suggestion
Slug messages	Separate messages for missing column, empty, invalid, several lines; "e.g." hint?	Yes, with the hint; missing column fails that sheet
Models on other environments	How to deploy model changes despite `mode="merge"`	Decide with the team before stage
UI colours	Blue / green / red banners for dry run vs real run; tinted Fail rows; AEM status colours; brand background (removed again)	Banners + Fail tint
Sling Job	Background uploads with progress	Skip until real bulk loads need it

Doubts still to verify

None of these block the demo; they need a check before a real rollout.

☐ **Excel date cells:** the guide says type dates as text; a real Excel date cell hasn't been tested yet.

☐ **`AssetManager` on Cloud:** saving the report works on the local SDK; confirm once on a Cloud dev environment.

☐ **Cloud request time limit:** time a real upload of about 500 rows to know how far 5,000 is from the limit.

☐ **Publishing options:** exact names of the tree-publish "only modified" options on your Cloud version.

☐ **Version purge:** report versions accumulate; confirm AEM's version purge settings on each environment.

☐ **Author permissions:** which groups can write in the fragment, imports and reports folders.

☐ **Model details from the photos:** `merchantLogo` as boolean, `preview` / `cache` as model fields (they may come from the Java model), and the label "offer Location" with a lowercase o were taken as shown; confirm with the source XML if available.

☐ **Header with a space and conversion errors** (`10001a`, `maybe`) haven't been run on the new models yet.

☐ **Local leftovers:** version 1.0 of the saved report has no label (from an early test), 22 old test workbooks sit in `/imports`, and empty folders remain under `offer-listing`.

Next steps

Suggested order after the demo: decide the deployment of models, close the data-safety gaps, then add convenience.

1. **Deployment readiness:** settle how model changes reach stage and prod; restrict the servlets to author; check author permissions; confirm `AssetManager` on a Cloud dev environment.
2. **Phase 2 validation:** duplicates (IDs and slugs), clearer slug messages, header check per sheet, then reference checks.
3. **Safety:** versions before UPDATE and before → after values in the dry run.
4. **Publishing:** start with Manage Publication; build the publish step if uploads become regular.
5. **Delete:** built 20 Oct; open: whether deleted IDs may be reused.
6. **Quality:** JUnit tests (deferred until the design is final), remaining tests (date cells, header with a space, conversion errors), local clean-up.
7. **Later / optional:** Sling Job for very large uploads, multi-value text fields, fragment metadata or variations, report walking every profile folder, UI colour polish.

Related documents

- CF Bulk Upload – Authoring Guide: for authors, first time to everyday use.
- CF Bulk Upload – Delete Challenges: decided and built 20 Oct.
- CF Bulk Upload – Empty Cell Clears Field: decided and built.
- Input format (spec): the contract the code implements.
- Test scenarios: every scenario with its workbook rows and verified result.