

ACS Commons Content Fragment Importer

Step-by-step operating guide and feature gaps

AEM as a Cloud Service | aemcloudproject | 22 September 2026

Chosen approach: Use the standard ACS Content Fragment Importer with six separate Excel files. No custom servlet, workflow, parser, or upload page. IDs, paths, and import order are prepared and maintained by the operator.

1. Confirm ACS and model readiness

- The project declares ACS AEM Commons 6.16.0, including the Cloud distribution. Verify the deployed instance has **Content Fragment Import** under Tools > ACS AEM Commons > Manage Controlled Processes > Start Process. [1]
- Deploy and enable all six models. Open each model and confirm the fields, types, and reference multiplicity match the workbook. The earlier package build succeeded; deployment and live authoring verification are separate steps.
- Before deployment, resolve the existing **/conf/aemcloudproject merge filter**. Merge does not replace old model definitions; renamed dialog nodes can leave duplicate fields. Use narrowly scoped model replacement, excluding those subtrees from the broad merge rule, and review existing-content migration. Do not replace the entire /conf tree. [4]
- Review prior property changes: merchantName to merchantName; Category name to Name; removal of child ID fields. Existing content and consumers may need migration.

2. Prepare folders and existing image assets

```
/content/dam/aemcloudproject/cfs/offer-listing  
[Offer Detail fragments are created here]  
categories/  
cards/  
offer-cta/  
merchants/  
venues/
```

Create the folders in Assets and configure the appropriate allowed-model policies. Upload required images and logos before importing. The original workbook uses /content/dam/sc/... image paths, while the updated picker roots use /content/dam/aemcloudproject; align these before testing.

Readiness check: Manually create one temporary fragment from each model in development and confirm its editor opens correctly.

Prepare the six Excel files

3. Keep one import table per workbook

Save six .xlsx files. Put the header row and import records on the **first worksheet**; other tabs are ignored. Use one complete row per fragment, no merged cells or continuation rows, and no extra explanation rows above the headers. [2]

4. Include the four technical columns

Column	Meaning
path	Destination folder, excluding the fragment name.
name	Stable asset node name, normally the slug.
title	Human-readable fragment title.
template	Complete path of the CF model.

Start with path, name, title, template, then the model fields. Populate all four technical values on every row; missing required values can cause records to be excluded. [2, 3]

DAM base: /content/dam/aemcloudproject/cfs/offer-listing

Model base: /conf/aemcloudproject/settings/dam/cfm/models

Workbook	path suffix	template suffix
01-Categories.xlsx	/categories	/category
02-Cards.xlsx	/cards	/offer-card
03-CTAs.xlsx	/offer-cta	/offer-cta
04-Venues.xlsx	/merchants/venues	/merchant-venue
05-Merchants.xlsx	/merchants	/merchant-details
06-Offers.xlsx	No suffix	/offer-detail

Append each suffix to the corresponding base above. Excel must contain the **complete paths**, not the words "DAM base" or "Model base".

5. Resolve the name / Name collision

Card and Venue have a model field named **name**; Category has **Name**. ACS 6.16.0 normalizes lookup keys to lowercase, so these share the mandatory technical name column. Do not add a second Name column for a separate display value.

With unchanged models: name = mastercard and title = Mastercard. The Card model name field also receives mastercard.
If a separate display value is required: rename the model properties and Excel headers to cardName, venueName, and categoryName. This is model configuration; the importer can remain standard.

Field checklist and cell formatting

6. Include the model data columns

Use property names only as headers. Remove notes such as "(fragment reference)" and "BasePath" from the original headers. In the unchanged-model approach, the technical name column also populates the name / Name model fields.

Model	Data columns after path, name, title, template
Offer Detail	offerID, offerSlug, offerTitle, offerSummary, offerDescription, offerImage, alternateText, hotPromo, offerStartDate, offerEndDate, offerURL, offerEmail, offerTnc, offerCategories, offerMerchants, offerCards, offerCta, offerOrigin, brCode, qrCode
Category	offerCategoryslug, categoryIconPath, parent Name is supplied by the technical name column.
Offer Card	offerCardSlug, logo, parent name is supplied by the technical name column.
Offer CTA	label, offerCtaSlug, url, deeplink
Merchant Venue	merchantvenueSlug, address, latitude, longitude, telephone name is supplied by the technical name column.
Merchant Details	merchantName, merchantSlug, merchantLogo, merchantVenues

7. Prepare values before upload

- **IDs:** Supply offerID yourself. Use stable allocated integers and check for duplicates; ACS does not generate business IDs or match by them.
- **Dates:** Use actual Excel date/time values, not ambiguous text such as 01/04/26. Confirm intended timezone and imported times with a small sample. [2]
- **Other types:** Use Boolean TRUE/FALSE for hotPromo, numeric cells for coordinates, and text cells for telephone numbers and codes.
- **Assets:** Use existing, complete DAM image paths. External CTA URLs and deeplinks go in their text fields.
- **Original sample cleanup:** Reconcile health-fitness with health-fitness-new; replace restrelax with an image path; combine the merchant continuation rows into one complete merchant record with a multiline venue-reference cell.

Updates require complete records. Keep all values that should remain on the fragment. A blank cell or omitted field is not a "leave unchanged" instruction in the inspected 6.16.0 importer. It can clear data or cause a conversion error.

References and import execution

8. Replace reference slugs with full fragment paths

The importer does not join workbooks by slug. A referenced child must be imported separately, and its complete asset path must be supplied in the parent record. For example, offerCards contains:

```
/content/dam/aemcloudproject/cfs/offer-listing/cards/mastercard
```

For several venues, put the following on separate lines **inside one merchantVenues cell**. Use actual in-cell line breaks, not the literal characters \n. [3]

```
/content/dam/aemcloudproject/cfs/offer-listing/merchants/venues/shilla-duty-free  
/content/dam/aemcloudproject/cfs/offer-listing/merchants/venues/shilla-duty-free-macau
```

The cell must contain exactly two complete paths, one per line.

Apply the same rule to offerCategories, offerMerchants, offerCards, and offerCta. Category/Card parent is a single reference and receives one complete path.

9. Import children before their parents

Run order: Categories > Cards > CTAs > Venues > Merchants > Offers.

Wait for each run to finish before starting its dependents. For category/card hierarchies, import new parent records in an earlier run. Worksheet row order does not guarantee processing order.

10. Execute each workbook in ACS

- Open **Tools > ACS AEM Commons > Manage Controlled Processes**.
- Click **Start Process** and choose **Content Fragment Import**.
- Upload the workbook; enable **Dry run** and **Detailed report**.
- Start the run, wait for completion, and inspect the process details and report.
- Start a new run with the same file. Disable Dry run and retain Detailed report.
- After completion, inspect the actual fragments in Assets, then proceed to the next workbook.

These controls are provided by ACS; no custom page or workflow is required. [1, 3]

Start small: Use one or two records per model in development. Dry run in the inspected version skips field writes; it is not a complete validation of values, references, permissions, or eventual import success.

Verify, update, and collect reports

11. Verify the first complete offer

- Open the created Offer Detail CF and check its ID, slug, title, dates, promotion flag, image, and all four reference collections.
- Open its referenced merchant and verify both venue links. Confirm multifields contain separate reference entries rather than one combined string.
- Check category/card parents, coordinates, telephone formatting, rich-text content, and CTA links. Confirm no accidental duplicate asset paths.

12. Test an update using the same asset address

Change one business field, such as offerSummary. Keep the path and name below unchanged and upload the complete Offers row again:

```
path = /content/dam/aemcloudproject/cfs/offer-listing
name = mount-faber-leisure-group
```

Operation	Stock importer behavior
Same path + name	Updates the existing fragment at that address.
Different path or name	Can create a new fragment, even if offerID is unchanged.
Change only offerID / offerSlug	Writes the field value; does not perform identity lookup or move the asset.
Change a reference list	Supply the complete desired list; no automatic append/merge.
Blank or missing model value	May clear content or fail conversion; not a partial-update contract.
Change technical title	In inspected 6.16.0, title is used during creation; existing fragment titles are not explicitly updated. Business fields such as offerTitle are separate.

13. Download and retain the process reports

Click the completed process. Use **View** to inspect its report or **Download (Excel)** to save it. Retain the input workbook, report, run date, and operator for each file. Check process errors as well as report totals. [1]

Reporting limitation in 6.16.0: The inspected library commits changes before checking whether to emit an Updated Fragment detail row. Successful updates can therefore be absent from that detail list. Verify changed CFs in Assets; the report is not a complete created-versus-modified audit.

14. Repeat safely

Keep the same path/name mapping for later updates. Check failures before rerunning: the importer commits individual work and does not roll back an entire multi-file batch. Importing does not constitute approval or publication; publish through your normal AEM process after verification.

What the stock-only approach misses

These are gaps against the original business story. They are not tasks implemented by the modified Excel format.

Original requirement	Gap / manual responsibility
Automatic business IDs	Enter and maintain IDs before import. No offerID generator.
Find by offerID or offerSlug	Matching uses path + name. Maintain a stable mapping yourself.
Link using child slugs	Resolve slugs into complete DAM fragment paths in Excel.
Create referenced children automatically	Import each child workbook separately; a reference alone creates no child CF.
One upload coordinates all entities	Only the first worksheet is read. The operator sequences the six files.
Cross-file/business validation	Check duplicate identities, unresolved relationships, parent cycles, date ranges, and asset paths before import. Model authoring validation is not a complete import validator.
Partial updates / reference merging	Prepare full records and full reference lists. No custom preserve, clear-marker, or append behavior.
Complete created/modified audit	Standard reporting is available, but detailed update entries have the 6.16.0 limitation described on page 5. No field-level change audit is supplied.
Atomic import and rollback	Partial success is possible. Inspect failures and recover/retry manually.
Automatic rename/move	Changing path/name can create a duplicate business entity. A controlled move is separate.
Separate business name and asset name	Current name / Name fields collide with technical name. Accept the slug value or rename the model fields.
Approval and automatic publication	Use the normal AEM approval/publication process separately.

Practical fit: Suitable when the team can maintain clean, complete spreadsheets and run imports in order. It does not meet automatic ID generation, business-key upsert, automatic child creation, or a complete modification audit without changing the scope or adding capabilities.

Operator checklist and references

Before each batch

- Models are deployed, enabled, and verified in AEM; no duplicate model fields.
- First sheet contains only clean headers and complete import rows.
- All path, name, title, and template cells are populated.
- IDs and slugs are supplied; path/name mappings are stable and unique.
- Image assets exist and reference paths match the imported children.
- Dates, Boolean values, coordinates, and text formats have been checked.
- Existing values and complete reference lists are retained on update rows.

During and after the batch

- Run the dry run and then a small real import before increasing volume.
- Complete child imports before merchants and offers; inspect each process for errors.
- Verify the first complete offer graph and a same-path update in Assets.
- Save inputs and reports; manually verify modified records where report details are absent.
- Review content and publish using the established AEM process.

Sources and verification basis

1. ACS Manage Controlled Processes - process launch and report download
2. ACS Spreadsheet API - first-sheet reading, required values, and dates
3. ACS Content Fragment Importer - input format and reference values
4. Apache Jackrabbit FileVault - package filter import modes
5. ACS importer source - implementation reference (master may change)

Version-specific behavior was checked by inspecting the locally cached `acs-aem-commons-bundle-6.16.0.jar`, including path/name matching, lowercase field lookups, field updates, dry run, title handling, and report timing. This was a code inspection, not an end-to-end import test on the running AEM instance. Recheck behavior if the deployed ACS version differs.